

Procedural Verifiability for Legal AI Outputs: An Audit Framework for Case Recommendation

Abstract

Legal AI systems can decide which legal authorities become visible before legal reasons are written. This article argues that this visibility power requires a distinct procedural-status test: legal AI outputs should qualify by their verifiable legal-material chain, not by model fluency, autonomy or retrieval accuracy alone. It introduces *normative material screening outputs*: outputs that filter, rank or summarize legal authorities without becoming legal sources or decision reasons. The category is formalized through output effect, system-role caps and a six-part audit vector covering source and corpus, query and version, authority hierarchy, ranking and counter-materials, contestability and adoption logging. The model is implemented in a provider-agnostic audit harness and evaluated through 230 scenario packets with 609 embedded records or outputs, 30 public source-text anchor checks, 50 raw transcript locator checks, 51,643 formal invariant checks, 185 metric-separation evaluations, 288 counterfactual gate ablations, threshold sensitivity, recoding and score-blinded codebook checks. The main result is an executable separation result: upstream authority coverage is insufficient for procedural qualification unless the output also preserves a reconstructable source chain, authority set, counter-material pathway and adoption record. Across 185 packets with upstream precision, recall and F1, recall above 0.8 predicted procedural qualification with precision 0.25; the audit gate cascade then removed the high-recall false positives. Across all 46 qualified packets, removing evidence, source-tag, authority, counter-material, review, role-cap or adoption gates caused 288/288 counterfactual downgrades. Thirty endpoint-matched public retrieval outputs remained reference information; ten raw GPT-5.5 outputs with full authority coverage were capped for lack of procedural-source binding; a 70-condition evidence ladder showed that the same model outputs change status only when source, contestability, logging and adoption gates change; and sixty adversarial repair variants were rejected despite high upstream recall. Public source-text verification anchored all 30 manifest support items, raw-output transcript verification anchored all 50 scenario locators, and invariant verification

exhausted the role-cap, evidence-packet, authority-set, counter-material, adoption and metric-separation gates.

Keywords: legal information retrieval; case recommendation; legal AI auditing; contestability; output evidence packets; normative material screening

1 Introduction

Legal AI systems increasingly decide what law becomes visible. A case recommender, generated research assistant or court-facing authority report may not decide a dispute, but it can set the field of authorities from which reasons are built. The central risk is not only hallucinated citations or poor retrieval; it is that an output quietly organizes the legal materials available for argument while being treated as neutral search.

This article’s claim is direct: procedural status should follow auditability, not intelligence. A fluent answer with weak sources remains reference information. A high-scoring ranked list with no counter-material pathway remains professional support. Only a reconstructable and contestable authority-screening chain can become procedural material. The missing category is *normative material screening*: outputs that filter, rank or summarize legal authorities without becoming legal sources or decision reasons.

The paper therefore asks three operational questions:

- RQ1: How should AI-generated legal outputs be classified according to their procedural effects?
- RQ2: What verifiability and contestability requirements are needed for each output class?
- RQ3: How can these requirements be operationalized and stress-tested across implementation, public-output, model-output, source-grounding, adversarial and recoding settings?

The resulting framework has two dimensions. The output-effect dimension classifies what the output does in the legal process. The system-role dimension classifies how the system enters that process. The two dimensions are then linked to verifiability, contestability and accountability requirements. The model is intentionally output-first: it asks when a legal AI output may move from internal reference to procedural material, and when it must be downgraded or withdrawn.

This article contributes a concept, a model and an artifact. The concept is normative material screening: the intermediate procedural role of AI-generated authority lists, ranked case sets and legal-material summaries. The model separates output effect from system role and links both to a six-part audit vector with status thresholds,

role caps and failure caps. The artifact is a provider-agnostic harness that implements the model and tests the paper’s central contrast: public search outputs can be source-reconstructable but procedurally underqualified; raw model outputs can identify the right authorities yet fail for lack of procedural source binding; source-supported repairs can qualify when source anchors, locators and counter-materials are validated; and adversarial repairs are rejected when source, locator, claim, link or counter-material gates fail.

AI-based case recommendation is the core case because it sits at the boundary between legal information retrieval and legal reasoning. Case recommendation systems are central to case-based reasoning and legal information retrieval, both longstanding areas of AI and Law research (Ashley, 1992; Rissland and Daniels, 1995; van Opijnen and Santos, 2017). They are also procedurally sensitive. In common-law precedent search, civil-law court-decision databases, administrative adjudication systems, arbitral award repositories and ODR platforms, ranking and omission can affect the legal materials available for argument. The framework is therefore jurisdiction-parameterized rather than jurisdiction-neutral: each legal system supplies its own authority hierarchy, while the verifiability problem remains shared.

2 Related Work and Gap

2.1 Legal information retrieval and case-based reasoning

Legal information retrieval and case-based reasoning both show why case recommendation cannot be evaluated as generic text ranking. Legal materials do not merely describe an external domain; statutes, precedents, administrative rules and court decisions may constitute or structure the legal domain itself. Legal retrieval must therefore account for hierarchy, temporality, authority, citations and professional use rather than reduce relevance to algorithmic similarity (Turtle, 1995; Rissland and Daniels, 1995; van Opijnen and Santos, 2017). Case-based reasoning reaches the same point from legal reasoning: cases matter through legally relevant similarities, differences, factors, values and precedential relations (Ashley, 1992; Bench-Capon and Sartor, 2003). The procedural risk is that a system may make some authorities salient and others invisible while presenting that ordering as ordinary search.

2.2 Benchmarks and legal retrieval models

Benchmarking work has substantially improved the technical evaluation of legal retrieval. COLIEE defines shared tasks for case-law retrieval, case-law entailment, statute retrieval and statutory entailment, with explicit datasets, runs and leader-

board metrics (Goebel et al., 2024). Neural retrieval systems such as BERT-PLI model paragraph-level interactions between a query case and candidate cases, addressing the length and structure of legal cases more directly than generic ad hoc retrieval models (Shao et al., 2020). Legal language and holding benchmarks such as LexGLUE and CaseHOLD extend evaluation to standardized legal-language understanding and legal holding selection (Chalkidis et al., 2022; Zheng et al., 2021). Recent retrieval-augmented legal benchmarks, including reasoning-focused legal retrieval and CLERC, move closer to professional research tasks by connecting retrieval to downstream legal analysis and citation support (Zheng et al., 2025; Hou et al., 2025). LegalBench further extends the evaluation agenda to a broader set of legal reasoning tasks for large language models (Guha et al., 2023).

These benchmarks are indispensable, but they answer a different question. COLIEE, LegalBench, CLERC and retrieval benchmarks ask whether a system retrieves, entails, cites or reasons correctly under task conditions. This article asks whether a particular output is procedurally eligible for use after it has been produced. Precision, recall, F1, mean reciprocal rank and entailment accuracy do not determine whether a ranked set may be attached to a court report, relied upon in a professional opinion, or incorporated into a decision reason. Procedural eligibility requires corpus records, authority labels, counter-authority handling, contestability and adoption logs that ordinary retrieval metrics do not measure.

2.3 Generated legal outputs and source verification

Generated legal interfaces make the procedural gap more important. A ranked list exposes a result set for inspection; a generated answer may compress selected materials into a fluent narrative and conceal omissions, conflicts or weak source support. Empirical work on legal hallucinations and commercial legal research tools shows that legal language fluency and retrieval augmentation cannot substitute for source verification (Dahl et al., 2024; Magesh et al., 2025). The audit task is therefore to distinguish source truth, retrieval relevance, authority status and procedural use: an output can be technically useful yet unfit for external procedural status, or partly opaque at the model level yet usable if the legal-material chain is reconstructable and contestable.

2.4 Auditability, explainability and contestability

The framework also draws on algorithmic accountability and AI auditing literature. Work on accountable algorithms emphasizes the need to review data, models and decision pathways rather than accept automated outputs as self-justifying (Kroll et al., 2017). AI auditing scholarship converts this concern into institutional practices for

examining systems before and after deployment (Raji et al., 2020; Mökander, 2023). Legal and social-science work on explanation and contestability further shows that explanation matters because affected persons need a meaningful opportunity to understand and challenge consequential decisions (Wachter et al., 2017; Binns et al., 2018).

Existing audit and explainability accounts, however, usually address automated decision systems in general. Due-process and public-sector accountability work shows why affected persons need a practical ability to challenge consequential automated support rather than receive a generic explanation after the fact (Citron and Pasquale, 2014; Veale et al., 2018). That insight has not yet been translated into a domain-specific account of legal AI outputs that structure the visibility and ranking of legal authorities without themselves deciding a case. This article fills that gap by treating case recommendation as normative material screening and by specifying the audit artefacts needed to qualify such outputs for procedural use.

3 Case Retrieval as Normative Material Screening

This article calls such outputs *normative material screening outputs*. The category is not a softer name for retrieval. It identifies an intermediate procedural object: an output that filters, ranks or summarizes legal materials that may later support legal argument or decision-making, without itself creating legal authority or deciding rights and duties. Examples include AI-generated lists of similar cases, recommended statutes, ranked precedent sets, extracted holdings and counter-authority reports.

The distinction also blocks two weak shortcuts. Retrieval metrics such as precision, recall, F1 and mean reciprocal rank are necessary for system development, but they do not say whether a recommendation may be cited in a procedural report, relied upon in a judicial memorandum or incorporated into a decision reason. At the same time, not every legal AI output that influences legal work is a decision-support reason. Many outputs shape the preparation of legal reasoning without becoming reasons for a decision. The middle layer is where most case recommendation systems actually operate.

4 A Two-Dimensional Output Qualification Model

The proposed framework starts from a simple premise: the normative qualification of a legal AI output depends on both what the output does and how the system enters the procedure. These are separate dimensions.

4.1 Output-effect dimension

The output-effect dimension classifies the legal effect claimed for the output:

The classification uses four criteria: whether the output is exposed outside the user’s internal workspace, whether a professional relies on it in advice or preparation, whether it filters the set of legal authorities available for argument, and whether an authorized decision-maker adopts it as part of a reasoned outcome. These criteria keep procedural status separate from model quality.

1. **Reference information.** The output supports internal search, summarization or drafting. It has no external procedural effect and is not presented as a basis for another person’s legal position.
2. **Professional support output.** The output assists a lawyer, expert, mediator or other professional in forming advice. The professional remains responsible for verification and communication.
3. **Normative material screening output.** The output filters, ranks, summarizes or highlights legal authorities or other normative materials. It may shape the materials available for argument, response or judicial consideration.
4. **Decision-support reason.** The output is adopted, summarized or incorporated by an authorized human or institution as part of the reasons for a legal decision, settlement recommendation, administrative action or arbitral outcome.
5. **Failure consequence: no external legal effect.** This is not a substantive output class. It is the cap applied when core verifiability, contestability or accountability requirements are absent.

4.2 System-role dimension

The system-role dimension determines how far an output may travel procedurally. It is not a second label pasted onto the same output. In the executable protocol, $R(o)$ is a role cap: a back-office tool may produce reference information, a disclosed assistance tool may produce professional support, an auditable procedural tool may produce normative material screening, and an authorized decision-support tool may produce decision-support reasons only after the adoption gate is satisfied. An unaccountable external disposition tool is capped at no external legal effect. The same ranked case list may therefore be private reference, professional support, court-facing screening, or a decision reason depending on the role and adoption path.

Table 1: Output status and procedural role

Output class	Typical use	Minimum procedural question	Audit artefact required
Reference information	Internal search, summary, first-pass drafting	Can the user verify the source before relying on it externally?	Source record and user-visible citation or document link.
Professional support output	Lawyer advice, expert preparation, mediator support	Can the professional explain the basis and limitations of the output to the affected person?	Source, query, output and limitation record.
Normative material screening output	Case recommendation, authority ranking, statute or precedent set selection	Can parties or reviewers reconstruct the corpus, version, ranking logic and omitted counter-materials?	Corpus manifest, ranking record, counter-material record and review log.
Decision-support reason	Output adopted in a judgment, administrative decision, arbitral award or ODR recommendation	Can the adoption path, human reasoning and contestation record be audited?	Adoption log, human reasons, contestation record and final responsibility node.
No external legal effect	Unaccountable or unverifiable output affecting rights, duties or procedural opportunities	Is any core requirement missing so that use must be blocked or downgraded?	Failure record stating the missing gate.

5 A Verifiability-Based Audit Model

The framework can be expressed as an audit model. Let o be a legal AI recommendation output generated in a particular institutional setting. Its procedural status $PS(o)$ is a function of output effect, system role and an audit vector:

$$PS(o) = g(E(o), R(o), \mathcal{A}(o)).$$

Here $E(o)$ denotes output effect, $R(o)$ denotes system role, and $\mathcal{A}(o)$ denotes a six-part audit vector:

$$\mathcal{A}(o) = (S, Q, H, K, T, L).$$

The six components are source and corpus verifiability (S), query and version traceability (Q), authority-hierarchy representation (H), ranking and counter-material verifiability (K), contestability support (T), and adoption logging with audit responsibility (L). Each component is scored on a three-point ordinal scale: 0 means absent, 1 means partially available, and 2 means procedurally reviewable. The model is not a predictive model of legal correctness. It is a status-allocation protocol: it determines the highest procedural status that the output may claim given the audit evidence available.

5.1 Formal definitions and invariants

Let the procedural statuses be ordered as

$$\mathcal{P} = \{p_0, p_1, p_2, p_3, p_4\},$$

where p_0 is no external legal effect, p_1 is reference information, p_2 is professional support output, p_3 is normative material screening output and p_4 is decision-support reason. The order $p_0 \prec p_1 \prec p_2 \prec p_3 \prec p_4$ expresses the maximum procedural status an output may claim. Let $\mathcal{A}(o) \in \{0, 1, 2\}^6$ be the audit vector and let $\Pi = (\tau_N, \tau_D, \Gamma, \Delta)$ be a status policy, where τ_N is the normative-screening threshold, τ_D is the decision-support threshold, Γ is the set of necessary gate predicates and Δ is the failure-flag cap.

System role is coded as $R(o) \in \mathcal{R}$, where $\mathcal{R}$ contains back-office tools, disclosed assistance tools, auditable procedural tools, authorized decision-support tools and un-

accountable external disposition tools. Define a role-cap function

$$c(R) = \begin{cases} p_1, & R = \text{back-office tool}, \\ p_2, & R = \text{disclosed assistance tool}, \\ p_3, & R = \text{auditable procedural tool}, \\ p_4, & R = \text{authorized decision-support tool}, \\ p_0, & R = \text{unaccountable external disposition tool}. \end{cases}$$

Output effect sets the claimed status; system role sets the maximum status available in the procedure. The operational rule then computes the highest procedurally available status from the audit vector, the role cap and the relevant adoption gate. For a candidate status p , define a gate predicate $m_p(\mathcal{A}, G_D, \Pi) \in \{0, 1\}$. The preliminary status is the greatest status whose audit gate is satisfied:

$$\sigma_\Pi(o) = \max_{\preceq} \{p \in \mathcal{P} : m_p(\mathcal{A}(o), G_D(o), \Pi) = 1\}.$$

Let $G_D = 1$ only when the review gate records completed review, authorized adoption, human authorization, jurisdiction assumptions, adoption reasons and contestation record. Failure flags then apply a cap operator $\kappa_F : \mathcal{P} \rightarrow \mathcal{P}$:

$$PS_\Pi(o) = \min_{\preceq} \{c(R(o)), \kappa_{F(o)}(\sigma_\Pi(o))\}.$$

Warning flags leave the status unchanged while recording a defect. Downgrade and suspension flags cap external status at reference information unless the missing artefact is restored. Withdrawal flags cap status at p_0 .

The protocol has three invariants. First, *gated monotonicity*: if $\mathcal{A}' \geq \mathcal{A}$ componentwise and R, G_D and F are unchanged, then $\sigma_\Pi(\mathcal{A}') \succeq \sigma_\Pi(\mathcal{A})$ only inside the same satisfied gate. Extra score cannot compensate for a missing source, counter-material or adoption gate. Second, *cap dominance*: for any output o , $PS_\Pi(o) \preceq \sigma_\Pi(o)$ and $PS_\Pi(o) \preceq c(R(o))$, while withdrawal flags force $PS_\Pi(o) = p_0$ regardless of upstream recall. Third, *metric non-equivalence*: retrieval metrics may inform K or issue-packet construction, but they are evidence variables, not status variables. These invariants are what make the model non-compensatory: legal fluency, high recall or strong ranking cannot purchase procedural status when a required audit artefact is absent.

Three consequences follow immediately. *Separation*: for any output with perfect upstream authority coverage but a non-procedural source tag, $PS_\Pi(o) \preceq p_1$ whenever it claims external screening status, because the source-attribution failure cap applies after g_0 ; missing output-source links are treated the same way. *Screening necessity*: an output cannot reach p_3 unless each audit dimension is at least partially present,

the output evidence packet is structurally present, and the counter-material gate is declared and reviewable; a high total score alone is insufficient. *Adoption necessity*: an output cannot reach p_4 unless $T = L = 2$ and $G_D = 1$, even if it already qualifies as screening. These are protocol properties, not empirical assumptions; the experiments below test whether the implementation enforces them on public, model-generated, repaired and adversarial packets.

5.2 Audit dimensions

This scoring scheme deliberately avoids requiring full model-parameter transparency in every case. A legal procedure may not need source code to decide whether a ranked case list can be contested. It does need enough information to reconstruct the legal-material chain and to test whether a relevant authority was omitted, misranked, outdated or mischaracterized.

5.3 Status thresholds

The audit vector determines the maximum status available to an output. A high aggregate score does not compensate for failure of a necessary gate. The model therefore uses both minimum gates and total-score bands.

The mapping function g is therefore explicit rather than discretionary. First, assign the highest status whose gates are satisfied:

$$g_0(\mathcal{A}) = \begin{cases} p_4, & S, Q, H, K, T, L \geq 1, \sum \mathcal{A} \geq 10, T = L = 2, G_D(o) = 1, \\ p_3, & S, Q, H, K, T, L \geq 1 \text{ and } \sum \mathcal{A} \geq 9, \\ p_2, & S, Q, L \geq 1, \\ p_1, & S, Q \geq 1, \\ p_0, & \text{otherwise.} \end{cases}$$

Second, apply system-role and failure caps. A disclosed assistance tool cannot exceed professional support even if the audit vector would otherwise satisfy normative screening; an auditable procedural tool cannot become a decision-support reason without authorized adoption. Withdrawal-level failures, such as unreconstructable sources, materially false generated summaries or irreversible external action without human authorization, set $PS(o)$ to no external legal effect. Suspension or downgrade flags, such as high-authority omission, invalid-authority recommendation, counter-material suppression, missing source-attribution tags, absent jurisdiction assumptions, review-gate failure, ranking drift or contestation failure, cap the output at reference status until

Table 2: Audit dimensions for case recommendation outputs

Dimension	Score 0	Score 1	Score 2
<i>S:</i> Source and corpus	Sources or corpus cannot be reconstructed.	Individual source links exist, but corpus scope or update status is incomplete.	Source links, corpus manifest, inclusion and exclusion rules, update date and document versions are recorded.
<i>Q:</i> Query and version	Query, filters or output version are unavailable.	Some query or version information is retained.	Query, filters, issue labels, source-collection version, workflow or interface version and output identifier are reconstructable.
<i>H:</i> Authority hierarchy	Output does not distinguish legal status.	Output marks broad categories such as binding or persuasive.	Output records jurisdiction, court level, legal force, validity, later treatment and relation to governing norms.
<i>K:</i> Ranking and counter-materials	Ranking is opaque and contrary materials are not surfaced.	Broad ranking factors or some contrary materials are shown.	Ranking/re-ranking factors, top- <i>k</i> set, omitted-material checks and counter-authority recall are reviewable.
<i>T:</i> Contestability	Affected persons cannot inspect or challenge the output.	A professional or internal reviewer can challenge the output.	Parties or reviewers can inspect, challenge, supplement and obtain a recorded response.
<i>L:</i> Adoption and responsibility	No adoption path or responsible actor is recorded.	Output and user action are logged.	Output ID, query, result list, adopted and rejected items, reasons, timestamp and responsible reviewer are logged.

Table 3: Status allocation rules

Candidate status	Necessary gates	Additional condition
Reference information	$S \geq 1$ and $Q \geq 1$	Output remains internal or user verifies sources before external reliance.
Professional support output	$S, Q, L \geq 1$	Professional user accepts verification responsibility and can explain limits.
Normative material screening output	$S, Q, H, K, T, L \geq 1$	Total score $\sum \mathcal{A}(o) \geq 9$ and counter-material handling is at least partially reviewable; declared complete-set omissions trigger caps.
Decision-support reason	$S, Q, H, K, T, L \geq 1$, $T = L = 2$ and authorized-adoption gate satisfied	Total score $\sum \mathcal{A}(o) \geq 10$ plus completed review, human authorization, reasons and recorded contestation.
No external legal effect	Any necessary gate for the claimed status is missing	Output is blocked, downgraded or retained only as internal reference.

the missing artefact is restored. Thus g combines the gate function g_0 with role and failure caps. If $K = 0$, the output cannot be treated as normative material screening because the ranking and counter-material treatment cannot be reviewed. If $T < 2$, $L < 2$ or $G_D(o) = 0$, the output cannot be a decision-support reason, even if retrieval accuracy is high. If the system purports to trigger an external disposition without legal authorization, review, contestability or accountability, the output has no external legal effect.

Operationally, an evaluator applies the model in six steps: identify the candidate procedural status claimed for the output; freeze the relevant issue, query, source collection, workflow version and output identifier; score each component of $\mathcal{A}(o)$ using the audit artefacts in Table 2; apply the necessary gates in Table 3; assign the highest available status or record the missing gate; and repeat the audit after material corpus, workflow or interface changes. This makes the framework usable as a deployment checklist, post-use audit and benchmark extension.

The thresholds use non-substitutable gates. The 0–2 scale is ordinal: 0 means absent, 1 means incomplete or internal, and 2 means reviewable by the relevant professional or procedural actor. The default $\tau_N = 9$ permits one limited weakness while still requiring every gate; decision-support status requires $T = L = 2$ because external reasons require both contestation and adoption records. The full validation suite re-

evaluates all 230 packets under thresholds 8–12. Thresholds 8–10 produced identical status allocations; only stricter thresholds generated demotions, and none generated promotions.

Table 4: Full-packet sensitivity of status allocation to threshold policy

τ_N	τ_D	Flips	Demotions	Status distribution
8	10	0	0	decision 12; screening 34; professional 23; reference 129; no effect 32
9	10	0	0	decision 12; screening 34; professional 23; reference 129; no effect 32
10	11	0	0	decision 12; screening 34; professional 23; reference 129; no effect 32
11	12	2	2	decision 12; screening 32; professional 25; reference 129; no effect 32
12	13	41	41	decision 0; screening 17; professional 52; reference 129; no effect 32

5.4 Jurisdictional parameterization

The audit dimensions are stable, but the materials satisfying H and K are local. For each setting j , evaluators instantiate authority hierarchy as H_j and counter-material handling as K_j before scoring.

5.5 Contestability and audit responsibility

Contestability requires a practical opportunity to inspect selected materials, identify missing or contrary materials, submit alternatives, request human review and receive a recorded response when the output is adopted. Adoption logging is distinct: it records whether, why and by whom the output was taken up after contestation was possible. Audit responsibility assigns each artefact to the actor that controls it: developers for capability and safety claims, deployers for corpus and interface records, professional users for issue framing and verification, and authorized adopters for decision uptake. The point is status evidence, not a general liability regime.

6 Operationalizing the Framework for Case Recommendation

AI-based case recommendation can be produced by many upstream architectures, but the audit framework does not depend on their internal mechanics. The legally relevant object is the output evidence packet: the source collection, versions, legal materials, output units, source locators, authority labels, counter-material links and adoption records that make the output reconstructable. Failures in these artefacts create different legal consequences. A stale source collection creates source and temporal risk. An ordering rule that privileges surface similarity may miss legally decisive distinctions.

Table 5: Jurisdictional profiles for H_j and K_j

Setting	H_j : authority hierarchy	K_j : counter-materials and limiting materials
Common law	Binding precedent, persuasive precedent, statutes, subsequent treatment, overruled or superseded authorities.	Distinguishing cases, contrary lines, limiting subsequent treatment, minority lines and overruled authorities.
Civil law	Constitutional norms, code provisions, special statutes, regulations, high-court decisions, lower-court decisions and doctrinal commentary.	Competing statutory interpretations, systematic conflicts, teleological limitations, contrary doctrinal views and later amendments.
Admin. adjudication	Statutes, binding regulations, agency rules, official guidance, adjudicative decisions and non-binding policy documents.	Conflicting agency interpretations, withdrawn guidance, higher-level statutory constraints and contrary adjudicative decisions.
Arbitration	Contract text, chosen law, mandatory law, institutional rules, prior awards and trade usage.	Contrary awards, mandatory-law limits, distinguishing clauses, procedural fairness objections and public-policy exceptions.
ODR	Platform rules, consumer statutes, terms of service, regulatory guidance and prior platform decisions.	Consumer-protection overrides, exception policies, appeal decisions and human-review overrides.

A generated summary may misstate holdings or flatten procedural posture. A user interface may display confirming cases while hiding contrary authorities.

The evaluation protocol should therefore separate upstream performance metrics from procedural checks. Precision, recall, F1, MRR and related metrics may help evaluate an upstream retrieval or recommendation component. They do not evaluate whether the output has preserved the legal hierarchy of authorities, displayed counter-materials, distinguished binding and persuasive materials, recorded update status, or produced a contestable adoption trail. These are evaluation requirements for procedural status.

The procedural metrics can be stated directly. *Authority coverage* asks whether known high-authority materials for the issue appear in the output set. *Counter-authority recall* asks whether legally relevant contrary or limiting materials are surfaced. *Hierarchy accuracy* asks whether the output correctly labels binding, precedential, persuasive, overruled or superseded materials. *Version traceability* asks whether the source collection and upstream run state that produced the output can be reconstructed. *Procedural-source coverage* asks whether source links are bound to official, tool-verified or otherwise procedurally usable source tags rather than merely named. *Adoption-log completeness* asks whether later human use of the output is observable. These metrics do not replace upstream benchmarks; they determine whether a benchmarked output can be used procedurally.

6.1 Baseline testing

Baseline testing should cover both technical and legal dimensions. A case recommendation system should be tested against known authority sets, especially high-authority or binding materials. Missing a binding, precedential, designated or otherwise high-authority material is different from missing one marginally similar low-level case. The former may trigger withdrawal from procedural use even if aggregate retrieval metrics remain acceptable.

Baseline testing should also record database coverage, update intervals, jurisdictional boundaries, language coverage, case hierarchy, citation treatment and exclusion rules. A system trained or indexed on a narrow corpus should not be evaluated as if it were a general legal research system. A system that excludes unpublished decisions, local court decisions or outdated but historically relevant authorities should make that exclusion visible to the user.

Table 6: Decision matrix for legal AI output qualification

Output use	Minimum verifiability	Minimum contestability	Allowed procedural status
Internal legal search	Source and version record; user can inspect retrieved materials	Internal user review	Reference information
Professional advice support	Source, query, output and limitations recorded	Client or supervising professional can request explanation when output materially shaped advice	Professional support output
Court or tribunal case recommendation	Corpus scope, update date, authority hierarchy, ranking factors, counter-authority handling and version record	Parties can inspect, challenge and supplement recommended materials	Normative material screening output
Decision reasoning support	Full adoption path, human confirmation, contrary-material treatment and audit trail	Parties can contest use; decision-maker gives reasons for adoption or rejection	Decision-support reason
Unaccountable external disposition	Missing legal authorization, review, contestability or accountability chain	Not satisfied	No external legal effect

6.2 Counter-material presentation

Case recommendation is procedurally dangerous when it confirms one line of reasoning without exposing plausible counter-materials. A useful legal recommendation system should therefore include counter-authority handling as a first-order evaluation requirement. The goal is not to force neutrality in the abstract. The goal is to preserve the adversarial and reason-giving structure of legal reasoning by ensuring that contrary cases, limiting cases, overruled authorities and minority lines are visible when they are legally relevant.

For audit purposes, counter-authority recall can be measured against a manually curated issue set:

$$CAR = \frac{|\text{retrieved counter-authorities} \cap \text{known counter-authorities}|}{|\text{known counter-authorities}|}.$$

The denominator is not the universe of all contrary law. It is the set of contrary or limiting materials identified for a defined issue by the relevant court, professional reviewer or benchmark curator. This makes the measure imperfect but usable.

This requirement is particularly important for generated legal outputs. A ranked list may visibly display multiple authorities. A generated summary may compress them into a single narrative and thereby hide disagreement. Evaluation should test whether the output cites supporting materials accurately, distinguishes holding from dicta, distinguishes binding from persuasive authorities, and signals unresolved conflict instead of hallucinating consensus. Empirical work on legal hallucinations shows why source verification cannot be left to model fluency (Dahl et al., 2024).

6.3 Output evidence packets

The audit layer proposed here sits above any particular search, retrieval, database, generation, agent or manual review architecture. The framework therefore does not inspect upstream implementation details. It requires an *output evidence packet*: a provider-agnostic record that connects the legal propositions or recommended materials in an output to source identifiers, versions, locators and review status.

A minimally reviewable evidence packet should record the issue presented to the system, source collection identifier, source version, output identifier, output units, source identifiers linked to each unit, paragraph or section locators, source-attribution tag, authority labels, counter-material links, human review gate, human adoption record and objection record. This requirement is related to model and dataset documentation work, but the unit of documentation is a procedurally relevant legal output rather than a model card or dataset sheet (Gebu et al., 2021; Mitchell et al., 2019). The upstream system may be a keyword search engine, legal database, generated-

output workflow, agent workflow or manual review process. The audit question is the same: can the legal-output chain be reconstructed, source-tagged, reviewed and contested?

Four structural packet metrics make this requirement operational. *Structural evidence fidelity* is the proportion of output-source links marked as supporting the output unit to which they are attached:

$$EF = \frac{|\text{output-source links that support the unit}|}{|\text{all output-source links}|}.$$

Evidence coverage is the proportion of legally material output units that have at least one reconstructable source locator:

$$EC = \frac{|\text{output units with source-locator support}|}{|\text{all legally material output units}|}.$$

Source-tag coverage measures output-source links carrying a verification-status tag:

$$STC = \frac{|\text{output-source links with source-attribution tags}|}{|\text{all output-source links}|}.$$

Procedural-source coverage is stricter:

$$PSC = \frac{|\text{links tagged official, tool-verified, user-verified or settled}|}{|\text{all output-source links}|}.$$

The tag may indicate whether the source was verified by a connected legal research tool, drawn from public metadata, supplied by the user, treated as settled background, or still requires pinpoint verification. These are packet-integrity metrics, not automated merits determinations: they test whether the output-source chain is structurally reviewable, while doctrinal support still requires professional or institutional review. For external normative screening, a tag that merely says “needs verification” records work still to be done. Low *EF*, *EC*, *STC* or *PSC* means that the output is not procedurally reconstructable even if an upstream system found relevant documents.

6.4 Review and adoption gates

Source reconstruction does not by itself authorize legal reliance. The gate record should state whether professional review is required and completed, what jurisdiction assumptions were used, which reliance level applies, whether any irreversible external action has human authorization, and what the human actor did with the output. An accurate draft remains below decision-support status until an authorized actor reviews it, adopts it and records reasons.

The minimum adoption-log schema is: output identifier, time, user role, issue,

filters, source-collection version, workflow version, displayed material set, displayed counter-materials, source-attribution tags, review-gate status, adopted item, rejected item, human reason, objection or supplement submitted, response to objection and final responsibility node. A system that cannot export these fields should remain below decision-support status.

7 Downgrade and Withdrawal Conditions

Procedural status is dynamic. Corpus changes, workflow updates, ranking drift, source-attribution gaps, review-gate failures or contestation failures can trigger failure caps after an output or system was initially approved. Table 5 operationalizes those caps: warnings correct isolated defects, downgrades cap status until a missing artefact is restored, suspensions block external use for an issue class, and withdrawals remove external use when the legal-material chain cannot be reconstructed or an irreversible action lacks human authorization.

Table 7: Downgrade severity levels

Severity	Trigger	Procedural consequence
Warning	Isolated error with reconstructable source, query and adoption record.	Correct the output and monitor the issue class.
Downgrade	Missing authority hierarchy, incomplete counter-material display, missing source tags, incomplete review gate or weak adoption logging.	Limit output to reference or professional support status until the missing artefact is restored.
Suspension	Repeated omission of high-authority materials, invalid authority recommendation or unexplained ranking drift in an issue class.	Suspend external procedural use for that domain or data source.
Withdrawal	Source, corpus, ranking or adoption chain cannot be reconstructed, the system generates materially false legal summaries, or an irreversible action lacks human authorization.	Withdraw from external legal use; retain only internal testing or research use.

8 Implementation and Evaluation

To test whether the framework is operational rather than merely classificatory, the rules above were implemented in a lightweight command-line audit harness. The implementation uses local JSON scenario files as the unit of evidence and produces both Markdown and JSON run artefacts. The companion artifact supplies the harness, scenario files, source snapshots, run artefacts and replication commands.

The harness implements the mapping in Section 4 directly. Each scenario declares a claimed status, system role or role-inference evidence, a base audit vector S, Q, H, K, T, L , a jurisdiction profile, upstream-performance metrics and expected outcomes. Authority sets, evidence packets, review gates and counter-material declarations are mandatory for outputs that claim or qualify for screening or decision-support status; missing fields trigger caps rather than silent upgrades. The runner does not read the expected outcome when computing status. It computes the allowed status, role cap, counter-authority recall, authority coverage, invalid-authority rate, structural evidence fidelity, evidence coverage, source-tag coverage, procedural-source-tag coverage and downgrade or withdrawal disposition. Separate commands vary status thresholds and re-score the same packets under strict and lenient coding policies.

The complete evaluation matrix contains 230 primary scenario packets with 609 embedded records or outputs, plus 30 public source-text anchor checks, 50 raw model-output transcript locator checks, 51,643 formal invariant checks, 185 metric-separation evaluations and 288 gate-ablation evaluations. The evidence is deliberately separated into five groups. Implementation and invariant tests check whether gates and caps execute as specified. Public-output audits test whether real legal search and listing outputs can be reconstructed and status-capped without model access. Model-output interventions test whether the same generated answers change status when source binding, counter-materials, contestability, logging and adoption gates change. Metric-separation tests compare upstream precision, recall and F1 thresholds with procedural status. Gate-ablation tests remove procedural conditions from every qualified packet. Robustness checks test threshold and coding fragility. The central empirical question is not whether every upstream system is accurate; it is whether the protocol separates retrieval success, source reconstructability and procedural qualification.

Derived robustness checks then vary thresholds, recode audit vectors under strict and lenient policies and run score-blinded codebook coding. These checks are not counted as independent field observations; they test whether the same packet evidence leads to stable status allocation.

Across ten stress-test scenarios, all rule outcomes were reproduced. The mean audit score was 8.30 and mean upstream recall was 0.84. Three scenarios had upstream recall

Table 8: Claim-to-evidence map

Claim	Evidence layer	What would falsify the claim
The protocol is executable	Stress scenarios, issue ablations and unit tests	Gates, caps or downgrade rules produce the wrong status.
Public outputs are auditable without model access	Metadata, public-system outputs and endpoint-matched retrieval	Visible records cannot be frozen, source-tagged or reconstructed.
Retrieval success is not procedural qualification	Public retrieval and raw GPT-5.5 outputs	High authority coverage alone upgrades the output to screening.
Retrieval metrics are not procedural status	Metric separation analysis	Precision, recall or F1 thresholds predict screening qualification with high precision.
Source-supported repair can upgrade an output	Repaired model outputs, evidence ladder and mixed-authority issue packets	Validated source, locator, counter-material and tag evidence still fails to qualify.
The validator rejects brittle repairs	Sixty adversarial repair variants	Locator, source, claim, link or counter-material attacks pass as screening.
Manifest support is externally anchorable	Public source-text anchor verification	Manifest support terms cannot be located in public source snapshots.
Raw model-output coding is transcript-anchored	Model-output transcript verification	Scenario locators cannot be found in the committed raw transcript.
Formal status properties are executable	Formal invariant verification	Scores alone can buy screening status without authority-set, evidence-packet, counter-material, review-gate, role-cap or adoption conditions.
Qualified outputs depend on their gates	Gate ablation analysis	Removing procedural gates from qualified packets leaves status unchanged.

Table 9: Full evaluation matrix

Layer	Units	Rule pass		What it checks
Stress scenarios	10	10/10		Status gates, failure caps, downgrade and withdrawal behavior.
Public meta-data	120 records in 6 files	6/6		Source reconstruction across six public legal-record sources.
Public-system outputs	60 records in 6 files	6/6		Real upstream legal-output reconstruction and professional-support caps.
Issue-specific public output/source packets	19 records in 3 files	3/3		Public issue-search outputs and one source-bound public-source packet tested against high-authority and counter-material requirements.
Endpoint-matched public retrieval benchmark	225 records in 30 files	30/30		Public case-law or known-item outputs tested against authority sets that the endpoint can return; mixed authority gaps are recorded separately.
Raw model outputs	10	10/10		High authority coverage without source binding remains procedurally capped.
Source-supported model-output repairs	10	10/10		Same model outputs qualify only after a validator checks manifest membership, locators, hashed source-support excerpts, source tags and issue-set coverage.
Model-output evidence ladder	70	70/70		Same model outputs move across reference, professional-support, screening, decision-support and no-effect status only when evidence and adoption gates change.
Adversarial source-support repairs	60	60/60		Negative controls test locator, claim, contradiction, source, link and counter-material attacks against the repair validator.
Public source-text anchors	30 checks	30/30		Support terms are checked against extracted public source text snapshots, reducing manifest-only validation.
Model-output transcript anchors	50 checks	50/50		Raw model-output scenario locators are checked against the committed transcript sections.
Formal invariant verification	51,643 checks	51,643/51,643		Exhaustive generated states test monotonicity, authority-set necessity, evidence-packet necessity, counter-material necessity, adoption gates, role caps, failure caps and metric non-equivalence.
Metric separation	185 evaluations	Recall-threshold precision 0.25; full gate precision 1.00		Tests whether upstream precision, recall and F1 predict procedural qualification.
Qualified-output gate ablations	288 ablations	288/288	22	Removes procedural gates from all qualified packets and checks that status falls below the corresponding level.
Mixed-	5	5/5		Source-bound statute, case and source

Table 10: Scenario-based harness results

Scenario type	Profile	Allowed status	Score	Up. recall	Audit result
Court authority report	Common law	Normative screening	10	0.83	Procedurally usable despite incomplete counter-material recall.
Civil-law statutory interpretation	Civil law	Normative screening	10	0.82	Uses competing interpretations and amendments as counter-materials.
Administrative guidance conflict	Admin.	Reference only	8	0.88	Strong upstream recall but suspended because counter-material and invalid-guidance flags fire.
Grounded output summary	Common law	Normative screening	11	0.90	Output-source mapping and structural evidence fidelity are complete.
High-coverage uncontestable output	Common law	No external effect	5	0.93	Strong upstream coverage does not cure missing contestability and weak evidence links.
Authorized ODR review	ODR	Decision-support reason	12	0.91	Permitted only because authorization, appeal, contestation and adoption logs are present.

above 0.80 but were procedurally blocked or downgraded. A high-recall output may still be unfit for external legal use when it omits a high-authority material, suppresses a relevant limiting material, recommends invalid authority, lacks output-source evidence or prevents affected persons from contesting the output.

8.1 Public legal-record reconstruction

The harness was also run on public legal-record snapshots. One layer sampled 20 public metadata records from each of six legal systems. A second layer froze the top ten records visible in six public legal retrieval or listing streams. These layers test reconstruction: whether visible records can be frozen, source-tagged, linked to locators and assigned a procedural cap by the audit layer. In both layers, all six scenario files were classified as professional support outputs because public record visibility establishes source reconstructability, not issue-specific ranking, counter-material completeness or party-facing contestation.

8.2 Issue-specific public output/source audit

A stricter public layer then froze two issue-specific public search outputs and one source-bound public-source packet, converting their visible records into evidence packets. The U.S. packet used a CourtListener issue search for agency deference after *Loper Bright*; the English-law packet used a National Archives issue search for mesothelioma causation after *Fairchild*; and the EU packet used selected Legislation.gov.uk and CURIA public source pages for GDPR Article 15 access-right materials. The layer contains 19 public records across three files.

The source-bound EU public-source packet qualified as normative material screening because it included the relevant high-authority materials, limiting provisions, source locators and reviewable output-source links. The U.S. and English issue-search outputs were capped at reference information because the visible returned records omitted high-authority or counter-material records for the defined issue. The procedural-status result is sharper than a relevance score: public legal records can be source-reconstructable yet fail normative-screening qualification when the output does not preserve the issue-specific authority and counter-material chain.

8.3 Public retrieval benchmark

To make that result less dependent on isolated examples, the harness also freezes 30 public retrieval outputs and scores them against endpoint-compatible authority sets. The benchmark covers five issue families: U.S. agency deference after *Loper Bright*, English mesothelioma causation after *Fairchild*, Canadian administrative-law

standard of review after *Vavilov*, German/EU right-to-be-forgotten review and EU GDPR Article 15 access rights. It uses CourtListener, the National Archives, Supreme Court of Canada Lexum search, OpenLegalData and CURIA public search or list endpoints. The benchmark contains 225 returned records. It is not a hand-built positive control: the returned top- k records are whatever the public endpoints exposed for the committed queries. Its gold sets are matched to endpoint scope, so case-law endpoints are not penalized for failing to return statutes; mixed statute/case screening is tested in the source-bound issue packets.

All 30 outputs were capped at reference information. Mean high-authority recall was 0.26 and mean counter-material recall was 0.00. Source-reconstructable public search output is therefore not the same as normative material screening. A system can expose real cases, stable URLs and rank order while still failing the procedural requirement that high-authority and limiting materials for a defined legal issue be made visible together.

8.4 Raw model-output pilot

The harness was also run on ten raw legal-authority recommendations generated by a single upstream model configuration, Codex GPT-5.5 with extra-high reasoning. The prompts covered the same three issue families as the positive-control packets: U.S. agency deference after *Loper Bright*, English mesothelioma causation after *Fairchild*, and EU GDPR Article 15 access rights. The model was not allowed to browse or use tools. This isolates the status question: whether fluent model text with strong authority coverage can itself claim normative screening status.

It cannot. In the coded raw-output packets, all ten outputs identified the expected high-authority and limiting materials in the issue families, producing mean upstream recall of 1.00 and counter-authority recall of 1.00. Yet every claimed normative-screening output was downgraded to reference information because the citations were marked as requiring verification rather than bound to official sources, corpus manifests, pinpoint locators and a contestation pathway. Their source-tag coverage was 1.00, but procedural-source coverage was 0.00.

The repair intervention then applies the opposite condition to the same model outputs. The visible authorities, issue families and model answers are held constant, but a validator checks each output-source link against an issue manifest, locator, hashed source-support excerpt, contradiction pattern, procedural source tag and issue-set coverage before assigning repair scores. A transcript verifier first checks that all 50 scenario locators appear in the ten committed raw-output sections. A separate public source-text layer checks the manifest support terms against extracted source snapshots; all 30 support items were externally anchored across 30 records with text snapshots. All ten

repaired variants qualified as normative material screening and none became decision-support reasons. This paired result is the paper’s central rule in executable form: authority coverage is not procedural status, but validator-backed source support can make authority coverage procedurally screenable when contestability and logging gates are satisfied.

The evidence-ladder intervention then varies seven conditions over the same ten outputs: raw unverified output, source-bound output without counter-material, source-bound output without contestability, source-bound output without logging, contestable screening, authorized decision-support adoption and unauthorized external action. The ladder produced the predicted status pattern: 30 reference-information outputs, 10 professional-support outputs, 10 normative-screening outputs, 10 decision-support reasons and 10 outputs with no external legal effect. This is the controlled result: the same generated answer changes procedural status only when the legal-material chain or adoption gate changes.

The same ten outputs were also converted into sixty adversarial repair variants. The attacks introduced locator mismatches, unsupported claims, contradiction-pattern claims, out-of-manifest sources, missing output-source links and counter-material omissions. All sixty adversarial variants were rejected despite mean upstream recall of 1.00: forty were capped at reference information, while twenty unsupported or contradictory source-support variants were withdrawn from external legal effect. The blocked-failure distribution across the full matrix shows why high-performing outputs failed: source-attribution gaps dominated, followed by summary distortion, counter-material suppression, authority omission and smaller numbers of contestability and invalid-authority failures.

Table 11: Raw model-output pilot

Issue family		N	Mean score	Upstream recall	PSC	Allowed status
U.S. agency deference after <i>Loper Bright</i>		3	9	1.00	0.00	Reference information
English mesothelioma causation after <i>Fairchild</i>		3	9	1.00	0.00	Reference information
EU GDPR Article 15 access rights		4	9	1.00	0.00	Reference information

To test mixed-authority normative material screening itself, the harness includes

five public source-screening packets: U.S. administrative law after *Loper Bright*, English-law mesothelioma causation after *Fairchild*, EU GDPR Article 15 access-right interpretation, Canadian administrative-law standard of review after *Vavilov*, and German/EU fundamental-rights review after the *Right to be Forgotten* decisions. Each packet defines high-authority materials, limiting or counter-materials, invalidated or historical treatments, source tags and output-to-source links at the level of individual output units. These are mixed statute/case/source construct tests, not a benchmark leaderboard: they ask whether a source-bound, issue-defined authority packet can obtain screening status under the protocol. All five qualified as normative material screening, not as decision-support reasons, because no authorized decision-maker adopted them as reasoned outcomes.

Table 12: Mixed-authority public source-screening packets

Issue packet	Profile	Authority set		Metrics	Status
U.S. agency deference after <i>Loper Bright</i>	Common law	H=3; invalid=1	C=2; treat- ment	Auth.=1.00; CAR=1.00; EF=1.00	Normative screening
English mesothelioma causation after <i>Fairchild</i>	Common law	H=3; invalid=1	C=2; treat- ment	Auth.=1.00; CAR=1.00; EF=1.00	Normative screening
EU GDPR Article 15 access rights	Civil law/EU	H=3; invalid=2	C=2; treat- ments	Auth.=1.00; CAR=1.00; EF=1.00	Normative screening
Canada standard of review after <i>Vavilov</i>	Common law	H=3; historical=1	C=2; treatment	Auth.=1.00; CAR=1.00; EF=1.00	Normative screening
Germany/EU right-to-be-forgotten review	Civil law/EU	H=3; C=2; cross- system=2	treat- ments	Auth.=1.00; CAR=1.00; EF=1.00	Normative screening

The ablation suite then removes one procedural condition at a time from those packets. Twenty ablations were generated: high-authority omission, counter-material suppression, unverified source tags and missing authorized-adoption gates for each issue family. Fifteen were capped at reference information through downgrade or suspension. The remaining five claimed decision-support status but were capped at normative screening because the packet lacked authorized institutional adoption. This is the crucial negative test: a packet that is otherwise strong on source links and authority coverage still loses status when the missing condition is procedural rather than semantic.

The metric-separation layer uses the 185 packets with complete upstream precision, recall and F1 scores to test whether conventional retrieval metrics predict procedural qualification. They do not. The point-biserial correlation between upstream recall and procedural qualification was 0.05. A recall threshold of 0.8 identified every procedurally qualified packet but also produced 136 false positives, yielding precision 0.25. The audit gate cascade then removes those high-recall false positives by adding score eligibility, source binding, review-role readiness and absence of failure caps. This result

is the empirical counterpart to the formal metric non-equivalence invariant: retrieval success may be necessary in many settings, but it is not a status-conferring condition.

The gate-ablation layer then tests gate necessity across the outputs that did qualify. For each of the 46 packets that reached normative screening or decision-support status, the harness generated counterfactual variants removing the evidence packet, replacing procedural source tags with verification-needed tags, omitting high-authority materials, omitting counter-materials, removing the review gate or imposing a professional-support role cap. For the 12 decision-support packets, it also removed the decision-adoption reason. All 288 ablations fell below the corresponding procedural level. The result matters because it shows that qualification is not carried by an isolated issue packet or a high total score; it depends on the procedural chain as a whole.

The final layers address coding uncertainty. First, all 230 committed scenario packets were re-scored under strict and lenient policies. Status allocation was stable for 218 of 230 packets; weighted status agreement was 1.00 for base versus strict and 0.98 for base versus lenient. Second, the same 230 packets were converted into score-blinded packet files that removed original scores, expected outcomes, source paths and manual failure flags. Two codebook-based coding passes then applied fresh scores to the packets. This is a codebook reproducibility check, not an external annotation study: treatment names and packet facts remain visible because coders must still read the legal-output evidence. The passes reached 0.99 exact and weighted inter-coder status agreement; against the base harness allocation, the weaker coder reached 0.95 exact agreement and 0.97 weighted agreement. The disagreements identify the fragile gates: decision-support adoption, deliberately thin stress packets and ablations near the screening threshold.

9 Implications and Limitations

The framework has three implications for legal AI evaluation. First, model performance is not procedural permission. A system may meet technical retrieval benchmarks while failing procedural qualification because the output cannot be reconstructed, source-tagged or contested. Second, human oversight is not magic. Oversight must leave review status, reasons, rejection records and adoption trails; a human signature without verifiable review should not upgrade an output’s status. Third, accountability should follow control. Developers, deployers, professional users and institutional adopters control different risks and should carry corresponding explanation duties.

The principal limitation is institutional realism rather than executable coverage. The current validation combines public search outputs, model-output interventions,

transcript anchors, public source-text anchors, adversarial repair attacks, metric-separation tests, gate-ablation tests, scenario tests, artifact checks and score-blinded codebook coding rather than a deployed-system field study with external human annotators. It shows how legal-output status can be audited above heterogeneous upstream systems and how robust the status rules are to source-binding attacks and coding uncertainty, while leaving adoption effects, user behavior and larger independent annotation for later work.

The next empirical step is therefore clear rather than open-ended. Deployed case recommendation systems and larger legal-output benchmarks should be scored for corpus verifiability, counter-material presentation, ranking traceability, adoption logging and contestation support, and those scores should be compared with retrieval metrics and user behavior. The question for future work is not only whether legal AI systems are accurate, but whether their outputs are legally usable.

10 Conclusion

Status follows auditability, not intelligence. That is the paper’s central rule. A case recommendation system may be technically sophisticated, fluent or empirically strong and still have no procedural status if its legal-material chain cannot be reconstructed and challenged. Conversely, an output does not need to be a legal source or autonomous decision-maker to matter procedurally. Once it structures the visibility and ranking of authorities, it occupies the middle category of normative material screening.

The proposed framework makes that category operational. It separates output effect from system role and links both to source, corpus, ranking, counter-material, contestability, adoption and audit requirements. It allows an output to remain internal reference information, become professional support, qualify as normative material screening, or enter a decision reason only through accountable human adoption. It also supplies downgrade and withdrawal rules when the evidentiary and procedural chain fails.

The contribution is therefore not another general claim that AI cannot be a judge or a legal source. It is a stricter evaluation logic for legal AI outputs: no source chain, no status; no counter-material pathway, no screening qualification; no adoption record, no decision reason.

Data and Code Availability

The companion artifact contains the audit harness, scenario packets, public snapshots, raw model-output transcript, validation scripts and generated Markdown/JSON re-

ports. The primary replication command is `pythonscripts/run_full_validation.py`; it reruns the full validation matrix and writes the aggregate report under `experiments/full_validation/results/`. The artifact is designed for deterministic review without LLM API access because public-source snapshots and raw model outputs are committed. Fresh public collection is available only through explicit refresh commands and is not required to reproduce the reported results.

References

- Ashley KD (1992) Case-based reasoning and its implications for legal expert systems. *Artificial Intelligence and Law* 1(2):113–208. <https://doi.org/10.1007/BF00114920>
- Bench-Capon TJM, Sartor G (2003) A model of legal reasoning with cases incorporating theories and values. *Artificial Intelligence* 150(1–2):97–143. [https://doi.org/10.1016/S0004-3702\(03\)00108-5](https://doi.org/10.1016/S0004-3702(03)00108-5)
- Binns R, Van Kleek M, Veale M, Lyngs U, Zhao J, Shadbolt N (2018) “It’s reducing a human being to a percentage”: Perceptions of justice in algorithmic decisions. In: *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*. ACM, pp 1–14. <https://doi.org/10.1145/3173574.3173951>
- Chalkidis I, Jana A, Hartung D, Bommarito M, Androutsopoulos I, Katz DM, Altras N (2022) LexGLUE: A benchmark dataset for legal language understanding in English. In: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*. ACL, pp 4310–4330. <https://doi.org/10.18653/v1/2022.acl-long.297>
- Citron DK, Pasquale F (2014) The scored society: Due process for automated predictions. *Washington Law Review* 89(1):1–33. <https://digitalcommons.law.uw.edu/wlr/vol89/iss1/2/>
- Dahl M, Magesh V, Suzgun M, Ho DE (2024) Large legal fictions: Profiling legal hallucinations in large language models. *Journal of Legal Analysis* 16(1):64–93. <https://doi.org/10.1093/jla/laae003>
- Dung PM (1995) On the acceptability of arguments and its fundamental role in non-monotonic reasoning, logic programming and n-person games. *Artificial Intelligence* 77(2):321–357. [https://doi.org/10.1016/0004-3702\(94\)00041-X](https://doi.org/10.1016/0004-3702(94)00041-X)

- Gebru T, Morgenstern J, Vecchione B, Vaughan JW, Wallach H, Daumé H III, Crawford K (2021) Datasheets for datasets. *Communications of the ACM* 64(12):86–92. <https://doi.org/10.1145/3458723>
- Goebel R, Kano Y, Kim MY, Rabelo J, Satoh K, Yoshioka M (2024) Overview and discussion of the Competition on Legal Information, Extraction/Entailment (COLIEE) 2023. *Review of Socionetwork Strategies* 18:27–47. <https://doi.org/10.1007/s12626-023-00152-0>
- Gordon TF, Prakken H, Walton D (2007) The Carneades model of argument and burden of proof. *Artificial Intelligence* 171(10–15):875–896. <https://doi.org/10.1016/j.artint.2007.04.010>
- Guha N, Nyarko J, Ho DE, Re C, Chilton A, Narayana A, Chohlas-Wood A, Peters A, Waldon B et al (2023) LegalBench: A collaboratively built benchmark for measuring legal reasoning in large language models. In: *Advances in Neural Information Processing Systems* 36. https://papers.nips.cc/paper_files/paper/2023/hash/89e44582fd28ddfea1ea4dcb0ebbf4b0-Abstract-Datasets_and_Benchmarks.html
- Hou AB, Weller O, Qin G, Yang E, Lawrie D, Holzenberger N, Blair-Stanek A, Van Durme B (2025) CLERC: A dataset for U.S. legal case retrieval and retrieval-augmented analysis generation. In: *Findings of the Association for Computational Linguistics: NAACL 2025*. ACL, pp 7913–7928. <https://doi.org/10.18653/v1/2025.findings-naacl.441>
- Kroll JA, Huey J, Barocas S, Felten EW, Reidenberg JR, Robinson DG, Yu H (2017) Accountable algorithms. *University of Pennsylvania Law Review* 165(3):633–705. https://scholarship.law.upenn.edu/penn_law_review/vol165/iss3/3/
- Magesh V, Surani F, Dahl M, Suzgun M, Manning CD, Ho DE (2025) Hallucination-free? Assessing the reliability of leading AI legal research tools. *Journal of Empirical Legal Studies* 22(2):216–242. <https://doi.org/10.1111/jels.12413>
- Mitchell M, Wu S, Zaldivar A, Barnes P, Vasserman L, Hutchinson B, Spitzer E, Raji ID et al (2019) Model cards for model reporting. In: *Proceedings of the Conference on Fairness, Accountability, and Transparency*. ACM, pp 220–229. <https://doi.org/10.1145/3287560.3287596>
- Mökander J (2023) Auditing of AI: Legal, ethical and technical approaches. *Digital Society* 2:49. <https://doi.org/10.1007/s44206-023-00074-y>

- Prakken H, Sartor G (2015) Law and logic: A review from an argumentation perspective. *Artificial Intelligence* 227:214–245. <https://doi.org/10.1016/j.artint.2015.06.005>
- Raji ID, Smart A, White RN, Mitchell M, Gebru T, Hutchinson B, Smith-Loud J, Theron D, Barnes P (2020) Closing the AI accountability gap: Defining an end-to-end framework for internal algorithmic auditing. In: *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*. ACM, pp 33–44. <https://doi.org/10.1145/3351095.3372873>
- Rissland EL, Daniels JJ (1995) A hybrid CBR-IR approach to legal information retrieval. In: *Proceedings of the 5th International Conference on Artificial Intelligence and Law*. ACM, pp 52–61. <https://doi.org/10.1145/222092.222125>
- Shao Y, Mao J, Liu Y, Ma W, Satoh K, Zhang M, Ma S (2020) BERT-PLI: Modeling paragraph-level interactions for legal case retrieval. In: *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence*. IJCAI, pp 3501–3507. <https://doi.org/10.24963/ijcai.2020/484>
- Turtle H (1995) Text retrieval in the legal world. *Artificial Intelligence and Law* 3:5–54. <https://doi.org/10.1007/BF00877694>
- van Opijnen M, Santos C (2017) On the concept of relevance in legal information retrieval. *Artificial Intelligence and Law* 25:65–87. <https://doi.org/10.1007/s10506-017-9195-8>
- Veale M, Van Kleek M, Binns R (2018) Fairness and accountability design needs for algorithmic support in high-stakes public sector decision-making. In: *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*. ACM, pp 1–14. <https://doi.org/10.1145/3173574.3174014>
- Wachter S, Mittelstadt B, Floridi L (2017) Why a right to explanation of automated decision-making does not exist in the General Data Protection Regulation. *International Data Privacy Law* 7(2):76–99. <https://doi.org/10.1093/idpl/ix005>
- Zheng L, Guha N, Anderson BR, Henderson P, Ho DE (2021) When does pretraining help? Assessing self-supervised learning for law and the CaseHOLD dataset of 53,000+ legal holdings. In: *Proceedings of the 18th International Conference on Artificial Intelligence and Law*. ACM, pp 159–168. <https://doi.org/10.1145/3462757.3466088>
- Zheng L, Guha N, Arifov J, Zhang S, Skreta M, Manning CD, Henderson P, Ho DE (2025) A reasoning-focused legal retrieval benchmark. In: *Proceedings of the*

4th ACM Symposium on Computer Science and Law. ACM, pp 169–193. <https://doi.org/10.1145/3709025.3712219>